

CELERY SCALING RESULTS FOR BATCH 1 TESTS

Purpose of this Document:

- To present result analysis of batch 1 tests conducted for Celery scaling.

Conditions of Tests:

- Digital Ocean VMs were used with the following configuration:
 - Codebase Server:
 - IP : 128.199.97.200
 - RAM : 1.96GB
 - Cores : 2
 - Operating System : CentOS 6.5
 - RabbitMQ Server :
 - IP : 128.199.116.114
 - RAM : 996.5 MB
 - Cores : 2
 - Operating System : CentOS 6.5
 - MySQL Server:
 - IP : 128.199.87.9
 - RAM : 996.5 MB
 - Cores : 2
 - Operating System : CentOS 6.5
- Celery version : 3.1.17
- RabbitMQ version : 3.4.4
- MySQL version : 5.6.14

Celery Configuration:

Directory Structure:

```
-- hod_app
  -- hod_app
    -- celery.py
    -- celeryconfig.py
    -- settings.py
```

- celery.py file created in inner hod_app folder.
 - Contains configuration for celery app instance.
 - Broker URL to be specified here.
 - Celeryconfig file linking also done here.
- Celeryconfig.py file created in inner hod_app folder.
 - All celery configuration and parameters to be defined here.
 - **Custom Exchanges : 7**

- `periodic_tasks` : Type --> Direct, routing_key --> 'periodic_tasks'
 - `sms_reminders` : Type --> Direct, routing_key --> 'sms_reminders'
 - `whatsapp_reminders` : Type --> Direct, routing_key --> 'whatsapp_reminders'
 - `phonecall_reminders` : Type --> Direct, routing_key --> 'phonecall_reminders'
 - `email_reminders` : Type --> Direct, routing_key --> 'email_reminders'
 - `generic` : Type --> direct, routing_key --> 'generic'
 - `billing` : Type --> direct, routing_key --> 'billing'
- **Custom Queues : 7**
 - `periodic_tasks` : Exchange --> periodic_tasks
 - `sms_reminders` : Exchange --> sms_reminders
 - `whatsapp_reminders` : Exchanges --> whatsapp_reminders
 - `email_reminders` : Exchange --> email_reminders
 - `phonecall_reminders` : Exchange --> email_reminders
 - `generic` : Exchange --> generic
 - `billing` : Exchange --> billing
- **Tasks Registered with celery : 12**
 - `billing.tasks.add_credits`
 - `billing.tasks.deactivate_expired_customer_purchases_and_realted_recipient_reminders`
 - `campaign.tasks.deactivate_expired_campaigns`
 - `customers.tasks.send_miss_call_verification_message`
 - `customers.tasks.send_sms_set_reminder_for_me`
 - `customers.tasks.sendmail_task`
 - `reminders.tasks.clean_voice_messages`
 - `reminders.tasks.make_phone_call_to_receipient`
 - `reminders.tasks.send_mail`
 - `reminders.tasks.send_reminders`
 - `reminders.tasks.send_sms_to_receipient`
 - `reminders.tasks.send_whatsapp_message_to_receipient`
- Custom Routes were defined to route tasks to their respective queues.
 - Tests were conducted using '`reminders.tasks.send_sms_to_receipient`'.
 - This calls a utility function '`send_sms`'.
 - URL of Project Sleuth used.
 - Request:


```
r = requests.post('http://128.199.87.42/receiver/api/post/request/',
data={'mode':'sms','token':'8tu30t284t0mc0m0bwl2'})
```

MySQL Seperation:

- MySQL was seperated from codebase server after initial test revealed primary memory bottleneck.
- Number of tasks fired : 1,00,000
- Remarks: Test could not be completed since MySQL crashed and celery was killed.
- Conclusion: Decision for MySQL seperation taken.

Units, Parameters and Definitions:

- Consumer Utilisation : ratio of publish rate and task completion rate.
- Publish rate, deliver rate, task completion rate measured per second.
- CPU Load measured as load per 1 minute average, 5 minute average, 15 minute average.
- Concurrency: The number of concurrent worker processes/threads/green threads executing tasks. If you're doing mostly I/O you can have more processes, but if mostly CPU-bound, try to keep it close to the number of CPUs on your machine. If not set, the number of CPUs/cores on the host will be used. Defaults to the number of available CPUs.
- CELERY_TASK_RESULT_EXPIRES : Time (in seconds, or a timedelta object) for when after stored task tombstones will be deleted. A built-in periodic task will delete the results after this time (celery.task.backend_cleanup). A value of none or 0 means results will never expire (depending on backend specifications). Default is to expire after 1 day.
- CELERY_IGNORE_RESULT : Whether to store the task return values or not (tombstones). If you still want to store errors, just not successful return values, you can set CELERY_STORE_ERRORS_EVEN_IF_IGNORED.
- CELERYD_MAX_TASKS_PER_CHILD : Maximum number of tasks a pool worker process can execute before it's replaced with a new one. Default is no limit.
- CELERY_DISABLE_RATE_LIMITS : Disable all rate limits, even if tasks has explicit rate limits set.
- CELERY_ACKS_LATE : Late ack means the task messages will be acknowledged after the task has been executed, not just before, which is the default behavior.

Total Test Cases : 7

◆ Test Case 1:

Initial Configuration:

- Codebase Server :
 - CPU Load : 0.02, 0.005, 0
 - Memory Consumption : 667 MB, 780 MB (with flower and workers)
- RabbitMQ server:
 - CPU Load: 0.08, 0.04, 0.01
 - Memory Consumption: 417 MB
- Celery Configuration:
 - Workers : 1
 - Concurrency : 20
 - CELERY_TASK_RESULT_EXPIRES = 3600
- Total tasks fired : 1,00,000

Server	CPU Load	Memory	Task Rate	Tasks	Publish/s	Deliver/s	Consumer Utilisation	Additional Information
Codebase	0.86,0.35,0.17	1.76 GB	17.58/s	33768	1000/s	-	-	-
RabbitMQ	0.95,0.90,0.14	658 MB	-	33768	1000/s	17.5/s	1%	-

◆ Test Case 2:

Initial Configuration:

- Codebase Server :
 - CPU Load : 0.01, 0.02, 0.0075
 - Memory Consumption : 368.55MB MB, 523 MB (with flower and workers)
- RabbitMQ server:
 - CPU Load: 0.06, 0.04, 0.01
 - Memory Consumption: 404 MB
- Celery Configuration:
 - Workers : 1
 - Concurrency : 30
 - CELERY_TASK_RESULT_EXPIRES = 1
 - CELERY_IGNORE_RESULT : True
 - CELERY_DISABLE_RATE_LIMITS = True
 - CELERY_ACKS_LATE = True
- Total Tasks fired : 1,00,000

Server	CPU Load	Memory	Task Rate	Tasks	Publish/s	Deliver/s	Consumer Utilisation	Additional Information
Codebase	1.04,0.46,0.23	1.75 BG	26/s	37022	998/s	-	-	Spike in memory consumption (550MB) observed. Could have been caused by Flower memory leak.
RabbitMQ	0.95,0.90,0.14	658 MB	-	33768	998/s	28/s	2%	

◆ Test Case 3:

Initial Configuration:

- Codebase Server : Similar to test casse 2
- RabbitMQ server : Similar to test case 2
- Celery Configuration: Similar to test case 2
- Total Tasks fired : 1,00,000
- Total tasks completed : 1,00,000
- This test was conducted to investigate memory spike.

Server	CPU Load	Memory	Task Rate	Tasks	Publish/s	Deliver/s	Consumer Utilisation	Additional Information
Codebase	1.04,0.46,0.23	1.75 BG	26/s	1,00,000	998/s	-	-	Memory spike not observed. Flower was not running.
RabbitMQ	0.95,0.90,0.14	658 MB	-	1,00,000	998/s	28/s	2%	

◆ Test Case 4:

Initial Configuration:

- Codebase Server :
 - CPU Load : 0.02, 0.03, 0.005
 - Memory Consumption: 212 MB, 255MB (flower, htop, top, 3 workers)
- RabbitMQ server :
 - CPU Load : 0.13, 0.07, 0.02
 - Memory Consumption: 404 MB
- Celery Configuration:
 - Workers : 3
 - Concurrency (per worker) : 10
 - CELERY_TASK_RESULT_EXPIRES = 1
 - CELERY_IGNORE_RESULT : True
 - CELERY_DISABLE_RATE_LIMITS = True
 - CELERY_ACKS_LATE = True
- Total Tasks fired : 50,000
- Total tasks completed : 50,000
- This test was conducted to investigate multiple worker scenario.

Server	CPU Load	Memory	Task Rate	Tasks	Publish/s	Deliver/s	Consumer Utilisation	Additional Information
Codebase	1.13, 0.68, 0.51	1.45 GB	26/s	50,000	1000/s	-	-	Memory spike not observed. Flower was not running.
RabbitMQ	0.77,0.27,0.09	542 MB	-	1,00,000	998/s	28/s	3%	

REMARKS :

- Multiple worker nodes (3 workers, 10 concurrency each) scenario performed equivalent (marginally better) as compared to single worker scenario (1 worker, 30 concurrency).
- Memory spike was not observed.
- Tasks were being equally distributed among the 3 worker nodes.
 - W1 : 16670 tasks
 - W2 : 16666 tasks
 - W3 : 16667 tasks
- Memory was not being released after completion of all tasks.
- Another test with similar configuration conducted without stopping workers. This was done to check any change in performance. No noticeable change in performance was observed.

◆ Test Case 5:

Initial Configuration:

- Codebase Server :
 - CPU Load : 0.035, 0.04, 0.16
 - Memory Consumption: 1.11GB
- RabbitMQ server :
 - CPU Load : 0.29, 0.09, 0.03
 - Memory Consumption: 410 MB
- Celery Configuration:
 - Workers : 3
 - Concurrency (per worker) : 17
 - CELERY_TASK_RESULT_EXPIRES = 1
 - CELERY_IGNORE_RESULT : True
 - CELERY_DISABLE_RATE_LIMITS = True
 - CELERY_ACKS_LATE = True
- Total Tasks fired : 50,000
- Total tasks completed : 50,000
- This test was conducted to investigate multiple worker scenario with higher concurrency.

Server	CPU Load	Memory	Task Rate	Tasks	Publish/s	Deliver/s	Consumer Utilisation	Additional Information
Codebase	2.54, 1.78, 1.33	1.56 GB	41.22/s	50,000	1000/s	-	-	Memory spike not observed. Flower was not running.
RabbitMQ	1.11, 0.51, 0.18	602 MB	-	50,000	1000/s	41/s	3%	

REMARKS :

- Total concurrency was 51, although task completion rate was 41.22/s.
- This shows bottleneck in task completion.
- CPU Load showed a maximum of 2.54 (1 min average) and was above 2 consistently throughout observation time.
- Hence CPU bottleneck for Dual core was reached.
- Memory spiker was not observed.

◆ Test Case 6:

Initial Configuration:

- Codebase Server :
 - CPU Load : 0.0, 0.01, 0.025
 - Memory Consumption: 203 MB, 724 (4 workers, shell_plus)
- RabbitMQ server :
 - CPU Load : 0.23, 0.08, 0.02
 - Memory Consumption: 392 MB
- Celery Configuration:
 - Workers : 4
 - Concurrency (per worker) : 12
 - CELERY_TASK_RESULT_EXPIRES = 1
 - CELERY_IGNORE_RESULT : True
 - CELERY_DISABLE_RATE_LIMITS = True
 - CELERY_ACKS_LATE = True
- Total Tasks fired : 50,000
- Total tasks completed : 50,000
- This test was conducted to investigate worker/concurrency limits of a 2GB, dual core DO VM.

Server	CPU Load	Memory	Task Rate	Tasks	Publish/s	Deliver/s	Consumer Utilisation	Additional Information
Codebase	3.26, 2.1, 1.57	1.73 GB	41.6/s	50,000	1000/s	-	-	-
RabbitMQ	0.94, 0.36, 0.17	610 MB	-	50,000	1000/s	41/s	4%	-

REMARKS :

- Total concurrency was 48, although task completion rate was 41.6/s.
- More tests with similar configuration were conducted and task completion rates were either equivalent or marginally better.
- A marginal increase in the task completion rate indicates that decrease in overall concurrency and moderate increase in number of workers, increases performance.
- Test shows that balance between number of worker pool processes and number of worker nodes must be reached through experiments and analysis.
- Memory spike was not observed.
- Next test was conducted without stopping workers, hence memory was not released.

◆ Test Case 7:

Initial Configuration:

- Codebase Server :
 - CPU Load : 0.57, 0.03, 0.53
 - Memory Consumption: 1.73 GB
- RabbitMQ server :
 - CPU Load : 0.23, 0.08, 0.02
 - Memory Consumption: 412 MB
- Celery Configuration:
 - Workers : 4
 - Concurrency (per worker) : 12
 - CELERY_TASK_RESULT_EXPIRES = 1
 - CELERY_IGNORE_RESULT : True
 - CELERY_DISABLE_RATE_LIMITS = True
 - CELERY_ACKS_LATE = True
 - CELERYD_MAX_TASKS_PER_CHILD = 200
- Total Tasks fired : 50,000
- Total tasks completed : 50,000
- This test was conducted to investigate effect of CELERYD_MAX_TASKS_PER_CHILD parameter.

Server	CPU Load	Memory	Task Rate	Tasks	Publish/s	Deliver/s	Consumer Utilisation	Additional Information
Codebase	4.02 2.57, 1.81	1.83 GB	40.6/s	50,000	1000/s	-	-	-
RabbitMQ	0.94, 0.36, 0.17	610 MB	-	50,000	1000/s	41/s	4%	-

REMARKS :

- Memory spike was not observed.
- Memory utilisation dipped as (almost) constant time intervals.
- CELERYD_MAX_TASKS_PER_CHILD could help prevent memory leak.

Conclusion:

- Target task completion rate to be achieved : 950 tasks per second.
- Limits of 2GB Dual core machine were observed.
- CPU processing limits proved to be a bottleneck.
- Maximum worker processes limit for a 2GB Dual core machine can be estimated at around 40 -45.
- Maximum task completion observed : 42/s
- Target Consumer utilisation must be 100%.
- MySQL separation had a positive effect on performance.
- Broker separation has a positive effect on performance.
- Further tests are being conducted on 8GB, quad core machines.
- Vertical limits of a single 8GB, 4 core machine will be estimated.
- Benchmarking of a 8GB machine will help in creating a generalised formula for task completion rate vs number of servers vs consumer utilisation.